


CS6462

Probabilistic and Explainable AI

Lesson 8

Principles of Probabilistic Programming

Bayesian Inference



Inference:

- use statistics to deduce properties about a probability distribution from data

Bayesian Inference:

- using Bayes' Theorem to do statistics inference
- Bayes' Theorem works not on events but on distributions:
 - Θ – set of parameters
 - *prior distribution* $P(\Theta)$ – distribution of our belief about the true value of Θ
 - *posterior distribution* $P(\Theta | \text{data})$ - distribution of our belief about Θ after we have taken the observed data into account
 - *likelihood distribution* $P(\text{data} | \Theta)$ - measures the degree to which data supports Θ


$$P(E|F) = \frac{P(F|E)*P(E)}{P(F)}$$

$$P(\Theta | \text{data}) = \frac{P(\text{data} | \Theta) * P(\Theta)}{P(\text{data})}; \quad L(\Theta | \text{data}) = P(\text{data} | \Theta)$$

normalizing factor

What kind of probability distributions should we use to model a probability?

Bayesian Inference (cont.)



Steps:


- Step 1. **[Prior]** Choose a *probability distribution function* to model your parameters Θ and the prior distribution $P(\Theta)$.
- Step 2. **[Likelihood]** Choose a *probability distribution function* for $P(\text{data} | \Theta)$. Basically you are modeling how **data** will look like given the parameters Θ .
- Step 3. **[Posterior]** Calculate the posterior distribution $P(\Theta | \text{data})$ and pick up the Θ that has the highest $P(\Theta | \text{data})$.

And the posterior becomes the new prior. Repeat step 3 as you get more data.

Probability Distribution Functions:

- *Normal distribution - has two parameters: mean μ and standard deviation σ*
- *Beta distribution*
- *Poisson distribution*

Bayesian Inference (cont.)



Normalizing factor $P(\text{data})$:

Why $P(\text{data})$ is important?

- the probability number that comes out is a normalizing factor
- recall: a necessary conditions for a probability distribution - the sum of probabilities of all possible outcomes S of an event is equal to 1, i.e., $P(S) = 1$
example: total probability of rolling a 1, 2, 3, 4, 5 or 6 on a die is equal to 1
- the normalizing factor makes sure that the resulting posterior distribution is a true probability distribution by ensuring that the sum of the distribution is equal to 1


$$P(\Theta | \text{data}) = \frac{P(\text{data} | \Theta) * P(\Theta)}{P(\text{data})}$$

normalizing factor

Ignoring $P(\text{data})$:

- could be ignored when the focus is on the peak of the distribution, regardless of whether the distribution is normalized or not

Bayesian Inference (cont.)



Posterior:


- The goal of Bayesian inference is to update our prior beliefs $P(\Theta)$ by taking into account **data** that we observe.
- If we assume that in any particular inference problem, data is fixed, we are interested in only the terms which are functions of Θ , i.e., ignore $P(\text{data})$:

$P(\Theta | \text{data}) \sim P(\text{data} | \Theta) * P(\Theta)$, or:

Posterior \sim **Likelihood** \times **Prior** (\sim posterior prob. distr., e.g., using Poisson distr.)

asymptotic (approximately equal)

- Our final beliefs about Θ combine both the relevant information we had a priori and the knowledge we gained a posteriori by observing data



Bayesian Inference (cont.)


Example: The probability of a certain medical test being positive is 90%, if a patient has disease . 1% of the population has the disease, and the test records a false positive 5% of the time. If you receive a positive test, what is your probability of having D?

$$\begin{array}{c}
 \text{Posterior probability} \quad \text{Likelihood function} \quad \text{Prior probability} \\
 P(D|+) = \frac{P(+|D) * P(D)}{P(+)} = \frac{P(+|D) * P(D)}{P(+|D) * P(D) + P(+|D^c) * P(D^c)}
 \end{array}$$

Result:

- *posterior probability* of having the disease **$P(\text{Disease}|+)$** given that the test was positive depends on the *prior probability* of the disease **$P(\text{Disease})$**
- **$P(+|D)=0.9, P(D)=0.01, P(+|D^c)=0.05, P(D|+) = ?$**
- Substituting in the numbers : **$P(D|+) = 0.15$**
- Final beliefs: **$\text{Posterior} \sim \text{Likelihood} \times \text{Prior} = 0.9 * 0.01 = 0.009$**

Bayes' Theorem & Probabilistic Programming



Bayesian Inference:

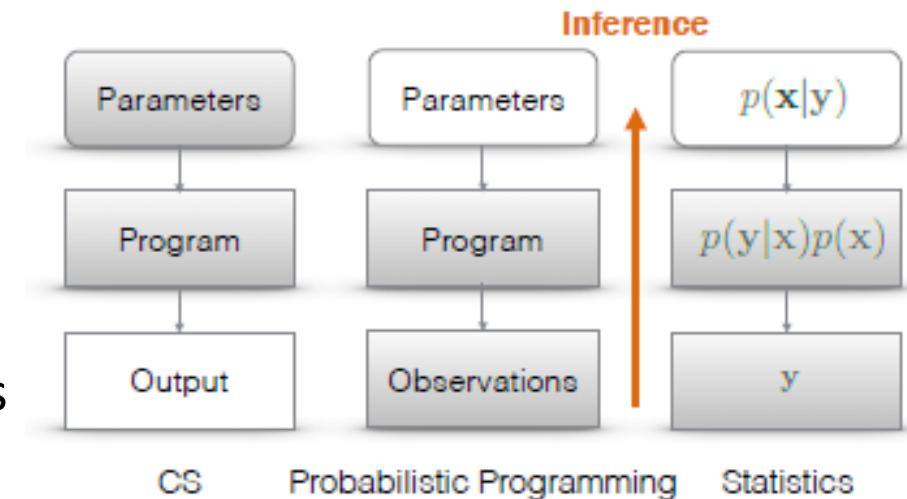
- Bayes' Theorem is used by statistical inference to update the probability for a hypothesis to cope with new information that has become available

Probabilistic Programming:


- about automating Bayesian inference
- syntax and semantics for languages that denote *conditional inference problems*
- formal semantics for building evaluators of models and applications from machine learning using the inference algorithms and theory from statistics

Example:

- \mathbf{y} – data (observations) output
- $p(\mathbf{y}|\mathbf{x})$, $p(\mathbf{x})$ – probabilistic models (data & parameters)
- $p(\mathbf{x}|\mathbf{y})$ – posterior distribution result of inference techniques



Probabilistic Programming

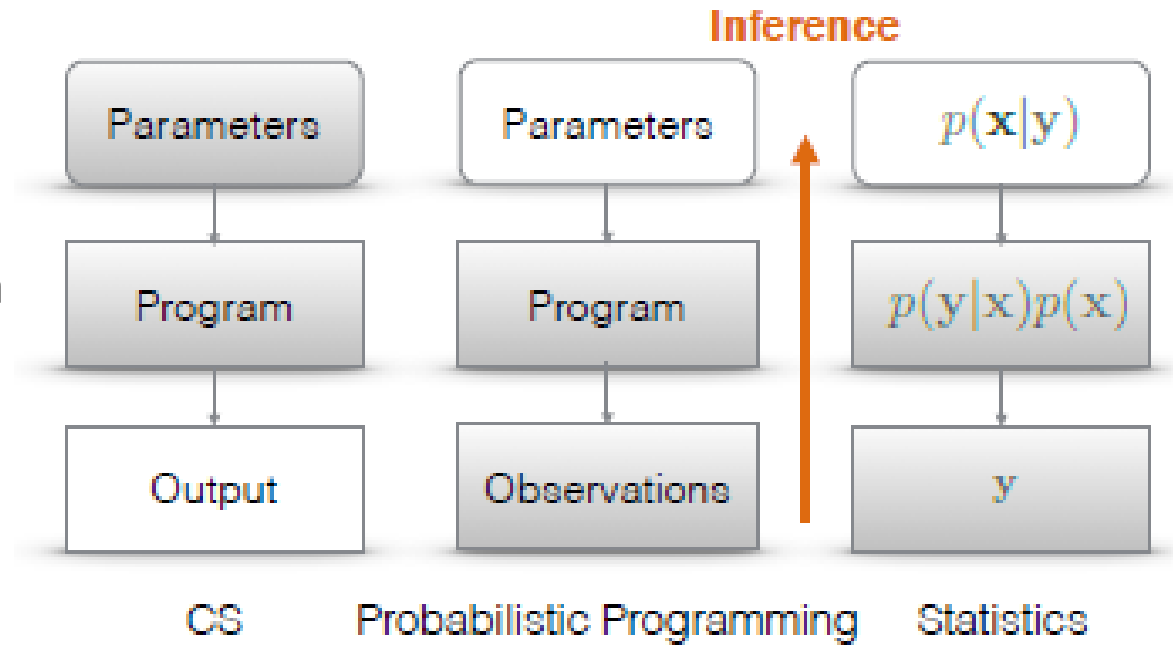


Example: disease test


- $y = \text{data}$: % of positive test records
- $x = \text{parameters}$ (having disease)
- $p(y|x)$, $p(x)$ (probabilistic models) - use Poisson distribution to build models for $P(+|D)$ and $P(D)$
- $p(x|y)$ (posterior distribution) –
Posterior \sim Likelihood \times Prior

Challenges:

- build the probabilistic models for Bayes' Theorem
- implement an algorithm following the theoretical model, so to computationally characterize the posterior distribution




Probabilistic Programming (cont.)



Features of probabilistic programs:

- probabilistic programs are functional or imperative programs with two added constructs [Gordon et al.]:
 - ability to draw values at random from distributions
 - ability to *condition* values of variables in a program via observations
- a probabilistic program simultaneously denotes a joint and conditional distributions:
 - indicates the conditioning
 - indicates what random variable values will be observed
- probabilistic programming languages support syntactic constructs for conditioning and evaluators that implement conditioning

Existing Probabilistic Languages



By research communities *:

- Anglican
- BLOG
- BayesDB
- Venture
- Probabilistic-C
- Church
- WebPPL
- CPProb
- Augur
- FOPPL
- Hakaru

* "An Introduction to Probabilistic Programming", book by J.W. van de Meent, B. Paige, H. Yang, F. Wood]

Probabilistic Program - Example



Example: reasoning about the bias of a coin

data - outcome heads or tails of one coin flip

model - Beta-Bernoulli model - a coin output and its bias are generated according to the model and then the coin flip outcome is observed and analyzed under this model

- y : data - heads a and tails b ; x - the bias of the coin $[0..1]$
- $p(x) \sim \text{Beta}(a, b)$ – prior probability distribution function
- $p(y|x) = p(a,b|x) \sim \text{Bernoulli}(x)$ – likelihood function
- $p(x|y) = p(x|a, b) \sim \text{Bernoulli}(x) * \text{Beta}(a, b)$ - posterior distribution

```
(let [prior (beta a b)
      x (sample prior)
      likelihood (bernoulli x)
      y 1]
  (observe likelihood y)
  x)
```

prior (beta a b) - function call that creates a prior distribution (heads and tails)


x (sample prior) - function call that creates a sample of x $[0..1]$

likelihood (bernoulli x) – the likelihood function

y 1 – y is assigned value 1 (heads)

(observe likelihood x if y) – posterior $p(y|likelihood)$

Summary



Bayesian inference:

- uses Bayes' Theorem to do statistical inference
- updates our prior beliefs by taking into account data that we observe
 - 1) ***Posterior*** \sim ***Likelihood*** \times ***Prior*** (\sim posterior prob. distr., e.g., using Poisson distr.)
 - 2) posterior becomes the new prior

Probabilistic Programming:

- about automating Bayesian inference
- syntax and semantics for languages that denote *conditional inference problems*
- formal semantics for building evaluators of models and applications from machine learning using the inference algorithms and theory from statistics

Challenges

Next Lesson – First-Order Probabilistic Programming Language (FOPPL)

Thank You!

Questions?